

## Processes - 2 The MUTEX Problem

Critical Section  
Basic Process Coordination  
Coordination Problems

## The critical section (CS) problem

Recall that a process is defined as a program that has started and not yet terminated.

### Definition:

Any code segment that accesses a shared data area is referred to as **Critical Section (CS)**!

A program is a **sequence of machine instructions** that are executed sequentially.

While we may assume that an individual machine instruction is **atomic**, a process may be interrupted after any one instruction

Shared data areas include:

- Memory Locations
- Files
- I/O devices
- Other **exclusive** resources

So what is so critical about this section?? Why do we care??

## ..CS example

Consider the following example from the literature:

```
cobegin
p1:  ...
    x = x+1;
    ...
||
p2:  ...
    x = x+1;
    ...
coend
```

Processes p1 and p2 will execute the following sequence:

$p_i: R_i = x; R_i = R_i + 1; x = R_i;$

Note:  $p_i$  can be interrupted after any one of the instructions.

This may result in different execution sequences

The corresponding machine instructions are:

1. load the value of  $x$  into internal register  $R$  ( $R=x$ )
2. increment  $R$  ( $R = R+1$ )
3. store the new value into memory location  $x$  ( $x = R$ )

## ...CS example

Consider the following 3 execution sequences and determine the value of  $x$ :

Sequence #1  
 $p1: R1 = x; R1 = R1 + 1; x = R1;$   
 $p2: R2 = x; R2 = R2 + 1; x = R2;$

Sequence #2  
 $p1: R1 = x; R1 = R1 + 1; x = R1;$   
 $p2: R2 = x; R2 = R2 + 1; x = R2;$

Sequence #3  
 $p1: R1 = x; R1 = R1 + 1; x = R1;$   
 $p2: R2 = x; R2 = R2 + 1; x = R2;$

How can we prevent from multiple processes being active in the CS??

## Towards a solution...

- Note that processes can be interrupted in the most inconvenient situations:
  - While evaluating the condition(s) in any programming construct:
    - while( cond )
    - for( ... )
    - until( .. )
    - if ( ... ) then
  - While inserting or removing an item from a data structure:
    - Tree
    - Linked list
    - Heap
- The solution to the CS problem is to allow at most one process to be active in the CS → Mutual Exclusion
- On way to accomplish this is to **disable interrupts** before entering the CS.
  - What if there are multiple CPUs?
  - What if processes have multiple (distinct) critical sections?
  - What if the CS is large?

## Mutual Exclusion

Any solution to guarantee **mutual exclusion** must limit the no. of processes in the CS to 1.

**Mutual Blocking** must be prevented!

We can/must distinguish the four types of blocking:

1. A process that is currently not executing in its CS must not prevent other processes to enter the CS (**progress**)
2. A process must not **repeatedly** enter its CS, thereby preventing other processes to enter the CS (**starvation**)
3. Two processes that are about to enter their CS must not block each other indefinitely (**deadlock**)
4. Two processes about to enter their CS must not repeatedly yield to each other indefinitely (**livelock**)

## Software solutions

The following are a number of "different" attempts to provide mutual exclusion for 2 processes! We need to evaluate each of them to understand what works and what doesn't!

Inspect each solution for all 4 types of blocking!

\\Mutual Exclusion is easy!!!

```
int in_CS = 0; // critical
cobegin
p1: while(1){
    while(!in_CS);
    in_CS = 1;
    DO-CS;
    in_CS = 0;
    prog1-outside CS }

p2: while(1){
    while(in_CS);
    in_CS = 1;
    DO-CS;
    in_CS = 0;
    prog1-outside CS }
coend
```



## ...from the literature ....

```
// Try #1
int turn = 1;
cobegin
p1: while(1) {
    while(turn == 2); //busy wait
    DO-CS;
    turn = 2;
    prog1 outside CS;
}
//
p2: while(1) {
    while(turn == 1); //busy wait
    DO-CS;
    turn = 1;
    prog2 outside CS;
}
```

Issues: Mutex? Blocking? Starvation?

DISCUSS!!

```
// Try #2
int c1 = 0, c2 = 0;
cobegin
p1: while(1) {
    c1 = 1;
    while(c2); //busy wait
    DO-CS;
    c1 = 0;
    prog1 outside CS;
}
//
p2: ... // analogous tp p1 //
```

Issues: Mutex? Blocking? Starvation?

DISCUSS!!

```
// Try #3
int c1 = 0, c2 = 0;

cobegin
p1: while(1) {
  c1 = 1;
  if (c2) c1 = 0; //busy wait
  else {
    DO-CS;
    c1 = 0;
    prog1 outside CS;
  }
}
//
```

```
p2: while(1) {
  c2 = 1;
  if (c1) c2 = 0; //busy wait
  else {
    DO-CS;
    c2 = 0;
    prog1 outside CS;
  }
}

coend;
```

Issues: Mutex? Blocking? Startvation?

DISCUSS!!

Note that we cannot predict the exact execution timing between p1 and p2!!

## ...a working solution

```
// Peterson Solution
int c1 = 0, c2 = 0, will_wait;
cobegin
p1: while(1) {
  c1 = 1;
  will_wait = 1;
  while(c2 && (will_wait == 1)); //loop
  DO-CS; c1 = 0;
  prog1 outside CS;
}

//
p2: while(1) {
  c2 = 1;
  will_wait = 2;
  while(c1 && (will_wait == 2)); //loop
  DO-CS; c2 = 0;
  prog1 outside CS;
}
coend;
```

- Process  $p_i$  sets flag  $c_i$  to indicate the intent to enter the CS.
- Variable *will\_wait* breaks possible race conditions.
- $p_i$  setting *will\_wait* to its *pid* announces to the other process that it is willing to wait if both processes happen to attempt to enter the CS at the same time.
- The solution guarantees mutual exclusion and prevents all forms of blocking.
- Formal Proof??

## ..notes on SW-based mutex

Software solutions have several **drawbacks**:

1. Solutions are often **difficult to understand and verify**.
1. **Extension** to more than 2 processes is generally difficult.
1. SW-based mutex solutions apply to competition problems, **not coordination/cooperation** among processes.
1. **Busy waiting** results in the utilization of the CPU by a waiting process without resulting in any computational progress.

- Proving the correctness of a mutex solution is not easy.
- In general, a good first approach is to **prove by contradiction**.
  - Assume mutex is violated
  - Show that reasoning leads to a contradiction.
- Sometimes it is easier to prove the contrapositive:
  - $A \rightarrow B$  (implication)
  - $\sim B \rightarrow \sim A$  (contrapositive)

## Test & Set → Semaphores

The crux of the problem to provide mutual exclusion is the unpredictable nature of processes. → when does a process get interrupted and is preempted??

One solution provided by most CPUs is the **test\_and\_set** instruction. There are various versions of TS in the literature.

The general format is:

test\_and\_set(x)

Semantics:

$x$  is a **boolean variable**, initialized to 0;

*test\_and\_set(x)* returns the value of  $x$  and sets  $x$  to 1.

```
boolean test_and_set(boolean x)
{
  test_and_set := x;
  x := 1;
  return;
}
```

*test\_and\_set is executed in a single instruction cycle and cannot be interrupted.*

## ...test\_and\_set example

note: there is an alternate form of TS in the book.

So, how does test\_and\_set help in providing mutual exclusion to a critical section?

Remember, TS is an **indivisible** or **atomic** operation provided by the CPU.

Hence, we can use TS to create a **spin lock** around a critical section!

Example test&set():

```
boolean lock := 0;
p:
  while(1){
    while test_and_set(lock);
    process CS
    lock := 0;
    remaining code outside CS;
  }
```

More on spin locks later!!

## Semaphores

A **semaphore S** is a data-structure that is maintained by the operating system.

Operation P and V on semaphores are **indivisible** or **atomic**.

How can the OS (i.e., software) provide atomic operations??

Guess!! Yes, YOU in the left corner.....

type semaphore...

```
typedef struct{
  int count;
  list_of_processes queue;
} semaphore;
```

```
P(s):    s.count --;
if s.count < 0 then block(s)
```

where block(s) places the process on s.queue and invokes the process scheduler.

```
V(s):    s.count ++;
if s.count <= 0 then wakeup(s)
```

where wakeup(s) removes a process from s.queue and places it into the ready list.

## Semaphore facts

- Semaphores were introduced by Dijkstra, 1968
- Operations P and V are acronyms for the Dutch words for:
  - P → to test → *proberen*
  - V → to signal → *verhogen*
- If several processes invoke P or V operations on the same semaphore, the operation will occur sequentially in arbitrary order.
- If more than one process is queued (inside a P operation) on the same semaphore, it is generally non-deterministic which process is selected upon the execution of V(s).
- However, this depends truly on the implementation of the semaphore and the scheduling discipline (or queuing discipline)
- Semaphores are maintained by the OS and are considered a system resource.