

Higher-Level Synchronization Mechanisms

AND-Synchronization
Sequencer and Event Count
Monitor

AND Synchronization

- As we have seen with the **Cigarette Smoker's Problem**, it is often difficult to coordinate access to resources by the means of simple semaphores.

- Why? What was the problem?

- Basically, we need to **test for the availability of multiple resources in one step**, i.e., tobacco and matches.

- We can do so with AND-Synch. or parallel semaphores.

We define new semaphore operations, **SP(s)** and **SV(s)**:

SP(s1,t1,d1;s2,t2,d2;.....;sn,tn,dn):
if $s1 \geq t1$ and $s2 \geq t2$ andand $sn \geq tn$ then for all i do $si := si - di$ else place process on queue associated with the first $si < ti$, establish program counter to restart the entire test.

SV(s1,d1;s2,d2;.....;sn,dn)
 $si := si + di$ for all i ; and wakeup all processes waiting on any of the si .

Strong Reader Preference

Shared Variables: NR is a semaphore initialized to the total number of readers, R; mx is a semaphore initialized to 1.

```
Reader
loop
    SP(NR,1,1);
    SP(mx,1,0);
    Perform read;
    SV(NR,1);
endloop;
```

```
Writer
loop
    SP(mx,1,1; NR, R, 0)
    Perform write;
    SV(mx,1);
endloop;
```

Writer Preference

Shared Variables: NW is a semaphore initialized to W, the total number of writers; NR is a semaphore initialized to R, the total number of readers; mx is a semaphore initialized to 1.

Writer	Reader
loop	loop
SP(NW,1,1);	SP(NR,1,1; NW,W,0);
SP(mx,1,1;NR,R,0);	Perform read;
Perform write;	SV(NR,1);
SV(mx,1; NW,1);	
endloop;	endloop;

Sequencer and Eventcounts

- An interesting mechanism for coordinating processes is the **Sequencer & Eventcount**.
- Consider a bakery or a bank with multiple sales persons/clerks.
- Multiple customers can enter the facility at the same time and must be organized such that they are served in order as soon as service is available.
- This problem has been solved by installing a ticket machine that issues tickets in numerical order.
- As soon as any of the servers become available, a counter that displays a number for the **Next Customer to be Served** is incremented.
- The same mechanism can be used to coordinate and synchronize processes.
- In OS, the stack or of tags by which customers are ordered is represented by an **Eventcount E**.
- The machine that issues tickets or tags to the customers is represented as a **Sequencer S**.

...Sequencer & Eventcount

- Only a single operation is defined on the sequencer S:
 - ticket(S)** :: issues a non-negative, increasing, and contiguous sequence of integers.
 - The **ticket(S)** operation corresponds to a newly arriving customer taking a unique numbered tag.
 - The eventcount E has 3 associated functions:
 - await(E,v);**
 - advance(E);**
 - read(E);**
- A sequence
- ```
v := ticket(s);
await(E,v)
```
- causes the process (customer) to wait until E reaches v.
- await(E,v)** suspends the calling process if  $E < v$ ; otherwise it allows the process to proceed.
- We may of course combine the two calls above to
- ```
await(E, ticket(s));
```

Sequencer & Eventcount

Formally, **await(E,v)** is defined as:

```
await(E,v):
if E < v then place the calling process in
the queue that is associated with E and
invoke the scheduler
```

advance(E) corresponds to an initiation of service: the eventcount value E is incremented and the next process/customer is admitted for service.

```
advance(s):
E := E + 1;
Wake up the process(es) waiting for E's
value to reach the current value just
obtained;
```

read(E) provides a means for inspecting the current value of E. This may be useful as a process may want to check how long it may have to wait.

Example: Solving the CS problem

```
E: eventcount /* initialized to 0 */
S: sequencer /* initialized to 0 */
```

```
await(E, ticket(S));
enter CS;
advance(E);
```

Example: Producer / Consumer

Shared Variables

```
var Pticket, Cticket: sequencer;
In, Out: eventcount;
buffer: array[0..N-1] of item;
```

!!! Remember, Pticket, Cticket, In, and Out are initialized to 0 !!!

```
Producer i:
var t: integer;
loop
Create new item;
t := ticket(Pticket)
await(In,t); /* one at a time */
await(Out, t-N+1);
buffer[t mod N] := item;
advance(In);
endloop;
```

```
Consumer k:
var u: integer;

loop
u := ticket(Cticket);
await(Out, u); /* one at a time */
await(In, u+1);
item := buffer[u mod N];
advance(Out);
consume item;
endloop
```

Monitor - an ADT


- On of the disadvantages of any of the synchronization mechanisms seen thus far is that they are very **low-level and hence prone to programming errors**.
- A **MONITOR** is a structure that contains data and functions.
- In other words, a monitor can be viewed as an **Abstract Data Type**, which facilitates the synchronization and coordination of access to objects.
- The advantages of the object oriented paradigm apply.
- The following are the characteristics of a monitor:
 - Mutual exclusive access** → only one process can be active inside the monitor.
 - Access to any of the encapsulated functions is only possible via the **monitor procedure**.
 - Use of **condition variables (CV)** → not really variables (no memory is associated with it).
 - Monitor semantics is based on **Dijkstra's Guarded Command**, e.g.: **[condition] → action**

..condition variables (CV)

- A condition variable is a name that is chosen by the programmer to represent a specific state or condition. For example:
 - condition: **CS_Open**
 may be used to represent the fact that the critical section is currently not used.
- Associated with each CV is a queue which can hold processes that are waiting for the condition represented by CV to become "true".
- The following functions are defined for CVs:
 - CV.wait** → causes the executing process to be suspended on the queue associated with CV.
 - CV.signal** → wakes up a process that waiting on CV, if one exists, otherwise no-op.
 - CV.queue** → true if queue is not empty, false if no process is waiting in the queue.
- We need to revisit the **semantics** after discussing the structure of a monitor.

Queues and more queues

- A monitor maintains many different queues:
 - Entry Queue → sequence processes to ensure mutual exclusive access.
 - CV Queues → store suspended processes that are waiting on the corresponding condition.
 - Urgent Queue → (to implement Hoare semantics) stores processes that have signaled others and are about to exit the monitor.



Monitor Examples...

A single resource allocator with Monitors

```
monitor: single_resource
{
    boolean: busy = false; /* initialization */
    condition: non_busy;

    acquire(){
        if busy
            non_busy.wait;
        busy = true;
    }

    release(){
        busy = false;
        non_busy.signal;
    }
}
```

An **extension** to CV.wait is **CV.wait(p)** where *p* is a priority number (the smaller *p*, the higher the priority)

Example: Shortest Job First (SJF)

```
monitor: SJN {
    condition x;
    boolean busy = false;

    acquire(int: time){
        if busy x.wait(time);
        busy = true;
    }

    release(){
        busy = false;
        x.signal;
    }
}
```

Different Semantics

- Recall, that only one process can be active inside the monitor at any given moment!
- We need to consider what happens if one process executes a CV.signal, thereby freeing a suspended process.
- This leads to two different CV.signal semantics:
 - Hoare Semantics
 - Mesa Semantics
- In the Hoare semantics
 - the process issuing the CV.signal is placed on the urgent queue if a process was waiting on the corresponding CV-queue.
 - it thereby yields the awakened process (only one process must be active).
 - Upon exiting the monitor, a process signals V(urgent) if there are processes suspended in the urgent queue.
 - If the urgent queue is empty, the exiting process will release access to the monitor.

...more semantics...

- In the Mesa Semantics:
 - the process issuing a CV.signal causes a waiting process to be placed on a ready queue.
 - CV.signal actually becomes a notify operation.
 - In this scheme, we cannot guarantee that a particular condition that was signaled to be true, is still true at the time the waiting process will execute.
 - Hence, a waiting process must re-evaluate the condition and suspend itself if it is found false.
- A process can easily re-evaluate the condition by using a while-construct instead of a simple if-construct to test the condition.
- ```
while(!B) c.wait;
```
- instead of the traditional
- ```
if(!B) c.wait;
```
- The Mesa semantics leads to a more efficient implementation and reduces the number of context switches.

Implementing a Monitor

- The following is an implementation of a Monitor ADT by means of semaphores:
- ```
P(mutex)
//only one process is active in the monitor
...
body of monitor procedures
...
if next.count > 0 V(next);
else V(mutex);

// where next is a semaphore that
// represents the urgent queue !!
```
- ```
CV.wait: →
CV.count ++;
if next.count > 0 V(next);
else V(mutex);
P(CV_sem);
CV.count --;
```
- ```
CV.signal: →
if CV.count > 0 {
 next.count ++;
 V(CV_sem);
 P(next);
 next.count --;
}
```

## Strong Reader Pref. with Monitor

The following is a solution to the strong readers priority using semaphores:

```
monitor: readers_writers;
int read_count;
boolean: busy = false;
condition: okto_read, okto_write;

{
 endread() {
 readcount --;
 if (readcount == 0) okto_write.signal;
 }

 startwrite() {
 if (readcount > 0 or busy)
 okto_write.wait;
 busy = true;
 }

 endwrite() {
 busy = false;
 if (okto_read.queue) okto_read.signal;
 else okto_write.signal;
 }

 startread() {
 if (busy) okto_read.wait;
 readcount ++;
 okto_read.signal;
 }
}
```

## ... some Monitor exercises

- Try to develop monitor solutions for the following problems:

- Dining Philosophers
- Bounded Buffer
- Cigarette Smokers

- Note: Monitor-based solution usually follow the same style →

- test and wait on CV are usually the first statements in a monitor procedure.
- signal or notify are usually the last actions that are performed.